Weibo Hot-Search Corpus — Data Summary

What we currently have: two collected datasets, their coverage, classification, and linkage to real posts

Snapshot as of 2026-09-30 · project root `~/weibo-corpus`

```
hot-search snapshots → daily aggregation → unique topics → 8-class LLM labelling → filter to 社会民生(S) + 国际涉外(N) → post retrieval (our Weibo account) → post URL + comment count + text excerpt [current end-point] → comment scraping (third party, not started)
```

52,857
topic-days collected (weibotop, 6.9 yr)

46,912
topics classified into 8 categories

18,019
topics linked to a real post (91%)

16,309
unique posts found, 59.5M comments located

1 Two datasets, not one superseding the other

We have run the full pipeline (classify → filter → find post) twice, against two different archives of Weibo's hot-search board:

Dataset	Source	Span	Topic-days	Fields
weibotop_2019_2026 (primary)	weibotop.cn, ~720 snapshots/day	2019-10-26 → 2026-09-19 (2,517/2,521 days)	52,857	title, peak hot, best rank, appearance count
lxw15337674_2025 (comparison)	GitHub archive, ~22 snapshots/day	2025-01-01 → 2025-12-31 (365 days)	64,858	+ Weibo's own category, description, read/discuss counts

The 2025 dataset ran first, to validate the method (classify → filter → find-post) before committing to the 7-year collection. Both are complete and usable; neither replaces the other. Full documentation for each lives in `data/sources/<name>/README.md`.

2 Primary dataset: weibotop, 2019–2026

Topic volume and category mix, by year

Year	Topic-days	S 社会民生	N 国际涉外	S+N	E 娱乐
2019 (67 days only)	1,407	25%	6%	31%	53%
2020	7,686	33%	18%	**52%**	36%
2021	7,665	33%	9%	42%	29%
2022	7,665	**36%**	10%	46%	**18%**
2023	7,665	30%	8%	37%	33%
2024	7,682	28%	9%	37%	31%
2025	7,585	27%	14%	40%	29%
2026 (to Sep 19)	5,374	29%	10%	40%	18%

2020 (COVID + early US-China friction) peaks on 国际涉外 at 18%, double most other years. 2022 (zero-COVID) peaks on 社会民生 at 36% and troughs on entertainment at 18%. The attention structure visibly shifts in 2020 and does not return to the 2019 baseline — composition itself is a finding, not just a sampling parameter.

8-class distribution, all years combined (46,912 labelled topics)

Code	Category	Count	Share
S	社会民生	14,602	31.1%
E	娱乐	13,552	28.9%
L	生活文化	5,864	12.5%
N	国际涉外	5,240	11.2%
T	体育	2,814	6.0%
C	科技商业	2,747	5.9%
P	官方宣传	1,523	3.2%
G	游戏电竞	570	1.2%

Classification is title-only (no description field available from this source), via DeepSeek V4 Flash. Validated against an independent Claude Sonnet annotator on a 600-item gold set: raw agreement 0.795, Cohen's κ 0.752 (8-class); 0.915 on the binary "is this S 社会民生" judgement that drives the filtering decision. Full validation in `data/sources/gold_standard/`.

3 Linkage to real posts — where things currently stand

From the 46,912 classified topics we selected **S (社会民生) + N (国际涉外)**, capped at the 10 highest-ranked per day, giving **19,814 topics**. Each was searched on `s.weibo.com` (our own logged-in account) to find the actual Weibo post people were discussing — not the hot-search phrase itself, but the specific post, its comment count, URL, and a text excerpt.

> **18,019 of 19,814 topics (91.0%) were successfully linked** to a real post. The search ran as one continuous session: 19,814 requests over 42.2 hours, never blocked. After deduplicating posts that multiple hot-search phrases pointed to (one event often spawns several hot-search terms across days), this is **16,309 distinct posts**.

By year

Year	Attempted	Linked	Hit rate	Median comments (top post)
2019	436	386	89%	2,139
2020	3,314	2,946	89%	**2,550**
2021	2,964	2,667	90%	1,133
2022	2,983	2,681	90%	1,543
2023	2,676	2,355	88%	910
2024	2,678	2,434	91%	1,097
2025	2,836	2,668	94%	866
2026	1,927	1,882	98%	**740**

Hit rate rises with recency (older hot-search phrases are more likely to have been reused for something unrelated since); median comment volume falls with recency (2020's median is 3.4x 2026's) — the two trends run in opposite directions and must both be reported.

Comparison dataset (lxw15337674_2025) — same pipeline, 2025 only, all four selected categories

Category	Attempted	Linked	Hit rate
S 社会民生	12,363	11,997	97%
N 国际涉外	5,966	5,759	97%
C 科技商业	4,080	3,595	88%
L 生活文化	8,231	7,291	89%

Higher hit rate than the weibotop run (93.5% overall) — this source's richer per-topic metadata and single-year scope likely both help. Confirms the linkage method generalises across both data sources rather than being an artefact of one.

4 What each linked record actually contains

For every linked topic, up to 5 candidate posts are stored (ranked by comment count, filtered to within ±7 days of the hot-search date): `p1` … `p5`, each with `comments, url, mid (Weibo internal ID), date, text` (first 400 characters of the post, captured from the same search response at no extra request cost). `p1` is the highest-comment post and is what "linked" above refers to.

5 What is not done yet

> **No comments have been scraped.** What exists is the location of each post (URL, ID) and its reported comment count — not the comment text itself. Comment scraping is a separate, later step, planned to be handled by a third party.

Volume exceeds the stated budget. Deduplicated across the 16,309 posts found, the sum of top-post comment counts is **~59.5 million** — roughly 4.6x the ~13 million comment budget. The excess is not evenly distributed: 8.2% of posts (those with >10,000 comments) account for 55% of the total, so raising a minimum-comment threshold barely helps (cutting to ⩾2,000 comments removes 64% of posts but only 11% of volume). The effective lever is instead a **per-post comment cap** — e.g. capping each post's scrape at 1,000 comments brings the total to ~11.3M (87% of budget) while still touching every one of the 18,019 linked topics. This threshold has not yet been finalised or applied to a delivery file.

All figures re-verified against the files in `data/sources/` on 2026-09-30. Per-source methodology, field definitions and known caveats: `data/sources/*/README.md` and `docs/DATA_SOURCES.md`. Sampling-design rationale and classification validation: `docs/SAMPLING_DESIGN.md`, `docs/report/methodology.html`.